

Message Queue Benchmarking Project

This project demonstrates and compares the performance characteristics of two popular message queue systems: Apache Kafka and RabbitMQ. The benchmark suite tests different throughput scenarios and measures latency and message processing capabilities.

Project Structure

```
.
├── benchmark/                # Benchmarking framework
│   ├── benchmark.py          # Main benchmark orchestration
│   ├── scenario.py           # Scenario definitions
│   ├── stats.py              # Statistics tracking
│   ├── utils.py              # Utility functions
│   ├── visualize_benchmark.py # Visualization tools
│   └── demo_logs/            # Benchmark results
├── message_queue/            # Message queue implementations
│   ├── interface.py          # Common interface
│   ├── factory.py            # Factory pattern implementation
│   ├── kafka.py              # Kafka implementation
│   └── rabbitmq.py           # RabbitMQ implementation
├── main.py                   # Entry point
└── README.md                 # This file
```

How We Demonstrate Kafka & RabbitMQ

The project uses a common interface to benchmark both Kafka and RabbitMQ under identical conditions:

1. **Common Interface:** Both implementations follow the same `MessageQueueBase` interface, ensuring fair comparisons.
2. **Docker Integration:** The system can run in both Docker and standalone environments, with appropriate broker address detection.
3. **Factory Pattern:** A message queue factory creates the appropriate implementation based on user selection.
4. **Command-line Selection:** Users can select which message queue to benchmark via command line arguments:

```
python main.py --queue kafka
python main.py --queue rabbitmq
```

Infrastructure Configuration

Docker Environment Setup

Our testing infrastructure is defined in `docker-compose.yml`, providing a consistent and isolated environment for benchmarking:

```
services:
  zookeeper:
    image: confluentinc/cp-zookeeper:latest
    environment:
      ZOOKEEPER_CLIENT_PORT: 2181
      ZOOKEEPER_TICK_TIME: 2000

  kafka:
    image: confluentinc/cp-kafka:latest
    depends_on:
      - zookeeper
    ports:
      - "9092:9092" # Internal
      - "9093:9093" # External
    environment:
      KAFKA_BROKER_ID: 1
      KAFKA_ZOOKEEPER_CONNECT: 'zookeeper:2181'
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka:9092,PLAINTEXT_HOST://localhost
      KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
      KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1

  rabbitmq:
    image: rabbitmq
    ports:
      - "5672:5672"
      - "15672:15672"
    environment:
      RABBITMQ_DEFAULT_USER: guest
      RABBITMQ_DEFAULT_PASS: guest
```

Kafka Configuration

Our Kafka setup is optimized for benchmarking with these key configurations:

- 1. Partitioning:** The Kafka topic uses 5 partitions, allowing for parallel message processing.

2. **Replication Factor:** Set to 1 for the benchmark environment (would be higher in production).
3. **Consumer Group:** Using a test consumer group to manage consumer offsets.
4. **Auto Offset Reset:** Set to "earliest" to ensure all messages are consumed during benchmarks.
5. **Dynamic Broker Detection:** Automatically detects whether to use the Docker container address or localhost.

Configuration snippet from `main.py` :

```
config = {  
    "kafka": {  
        "topic": "live_demo",  
        "partition": 5,  
        "replication_factor": 1,  
        "producer": {"bootstrap.servers": kafka_broker},  
        "consumer": {  
            "bootstrap.servers": kafka_broker,  
            "group.id": "test_group",  
            "auto.offset.reset": "earliest"  
        }  
    }  
}
```

RabbitMQ Configuration

Our RabbitMQ setup is configured for optimal performance comparison:

1. **Multiple Queues:** Using 5 queues to match Kafka's partition count for fair comparison.
2. **Authentication:** Using default guest credentials for simplicity.
3. **Connection Parameters:** Configured with appropriate heartbeat and timeout settings.
4. **Queue Properties:** Standard durability settings for persistence.

Configuration snippet from `main.py` :

```
config = {  
    "rabbitmq": {  
        'host': 'localhost',  
        'port': 5672,  
        'username': 'guest',  
        'password': 'guest',  
    }  
}
```

```
    'queue': 'live_demo',  
    'num_queues': 5,  
  }  
}
```

Additional parameters in the implementation:

- Heartbeat: 60 seconds
- Connection attempts: 3
- Retry delay: 5 seconds
- Socket timeout: 10 seconds

Benchmark Scenarios

Our benchmark suite includes five carefully designed scenarios that progressively test different aspects of message queue performance:

1. Low Throughput Scenario

Configuration:

- 500 total messages
- Sent in batches of 10 messages
- Tests baseline performance with minimal load

Implementation Details:

- Messages are produced in small batches with minimal delay
- All consumers are active throughout the test
- Measures both per-message latency and end-to-end processing time
- Represents typical performance for lightweight applications with occasional messaging needs

Observations:

- Interestingly, RabbitMQ outperforms Kafka in this scenario (0.26ms vs 6.18ms latency)
- RabbitMQ's simpler architecture provides advantages for low-volume scenarios
- Kafka's overhead for partitioning and consumer group management shows in these results

2. Medium Throughput Scenario

Configuration:

- 10,000 total messages
- Sent in batches of 50 messages
- Small delay (10ms) between batches
- Tests performance under moderate, sustained load

Implementation Details:

- Simulates a typical production application with moderate throughput requirements
- Consumer processing speed becomes more important than in the low throughput scenario
- Demonstrates how message batching affects overall system performance

Observations:

- Performance between Kafka and RabbitMQ becomes more comparable (0.44ms vs 0.47ms)
- End-to-end latency remains similar (4.43s for Kafka vs 4.71s for RabbitMQ)
- This similarity suggests that for medium workloads, architectural differences have minimal impact

3. High Throughput Scenario

Configuration:

- 100,000 total messages
- Sent in batches of 100 messages
- Minimal delay (10ms) between batches
- Tests performance under heavy, continuous load

Implementation Details:

- Represents high-volume production systems like log aggregation or event streaming
- Pushes consumer processing to near capacity
- Message ordering and system stability under load become critical factors

Observations:

- Kafka now shows significant latency advantages over RabbitMQ (0.20ms vs 0.37ms)
- End-to-end processing time difference becomes substantial (20.20s vs 36.76s)
- Demonstrates Kafka's architectural advantages for high-throughput scenarios
- RabbitMQ shows signs of increased resource utilization but maintains stable processing

4. Consumer Disconnect Scenario

Configuration:

- 100,000 total messages
- Sent in batches of 100 messages
- One consumer randomly disconnected halfway through processing
- Tests resilience and recovery capabilities

Implementation Details:

- Simulates real-world system failures or maintenance scenarios
- At message #50,000, a randomly selected consumer is shut down
- Remaining consumers must compensate for the lost processing capacity
- Implementation in `scenario.py` :

```
if index == 500:
    consumer_index = random.randint(0, self.benchmark_utils.message_queue.consumers - 1)
    self.benchmark_utils.message_queue.stop_consumer(consumer_index)
```

Observations:

- Detailed in the "Understanding the Consumer Disconnect Scenario" section
- Reveals fundamental architectural differences in failure handling
- Kafka prioritizes complete message processing at the cost of increased latency
- RabbitMQ maintains performance but may not process all messages

5. Extreme Throughput Scenario

Configuration:

- 1,000,000 total messages
- Sent in batches of 100 messages
- Minimal delay between batches
- Tests absolute performance limits

Implementation Details:

- Pushes both systems to their practical limits
- Represents extreme use cases like high-frequency trading or IoT sensor networks
- Measures system stability and consistency under extreme pressure
- Helps identify potential bottlenecks in both implementations

Observations:

- Kafka maintains impressively low latency per message (0.17ms)
- RabbitMQ's latency remains consistent with other scenarios (0.36ms)

- The end-to-end processing time difference becomes dramatic (167.54s vs 363.23s)
- Kafka processes all messages more efficiently, showing better scaling characteristics
- RabbitMQ successfully processes all messages but with significantly higher total processing time

Results

The benchmark captures two key metrics:

1. **Average Latency per Message:** The average time it takes to process a single message
2. **End-to-End Latency:** The total time to process all messages in a scenario

Visualization of Results

The following visualizations illustrate the performance characteristics of Kafka and RabbitMQ across all benchmark scenarios:

Latency Comparison

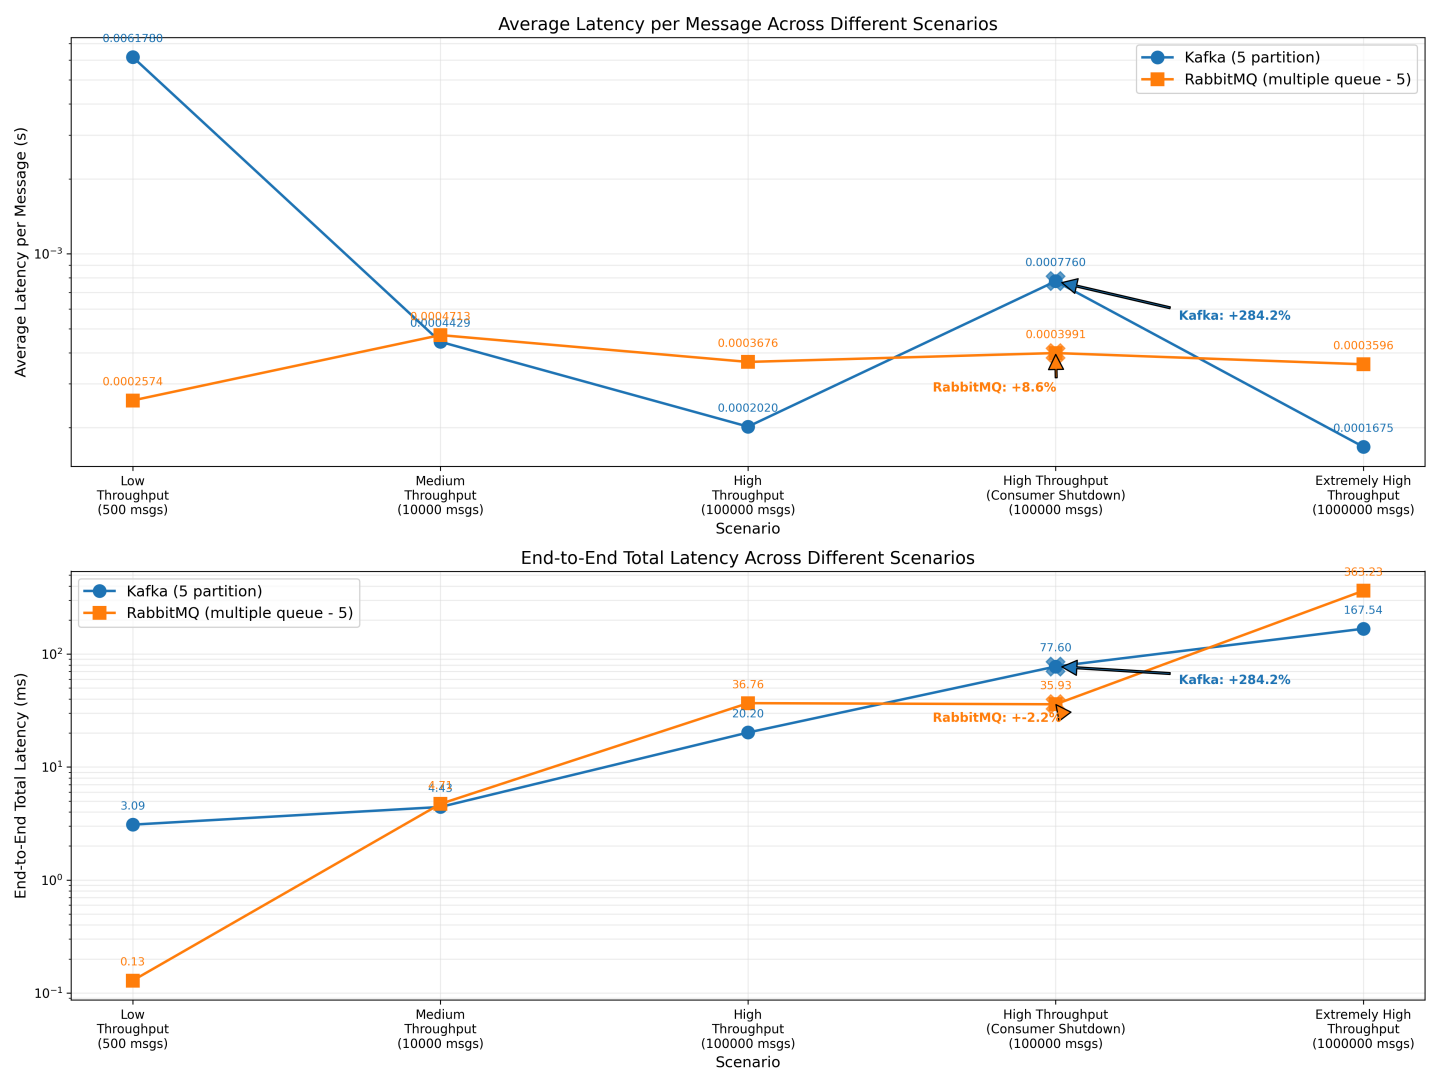


Figure 1: Comparison of average message latency and end-to-end latency for Kafka and

RabbitMQ across different scenarios. Note the logarithmic scale for latency values.

The chart shows both the average latency per message (top) and the end-to-end total latency (bottom) across all scenarios. Key observations:

- Kafka generally maintains lower per-message latency except in the low throughput scenario
- The consumer disconnect scenario shows a notable spike in Kafka's latency
- RabbitMQ maintains relatively consistent latency across scenarios
- End-to-end latency differences become more pronounced as message volume increases

Message Consumption Comparison

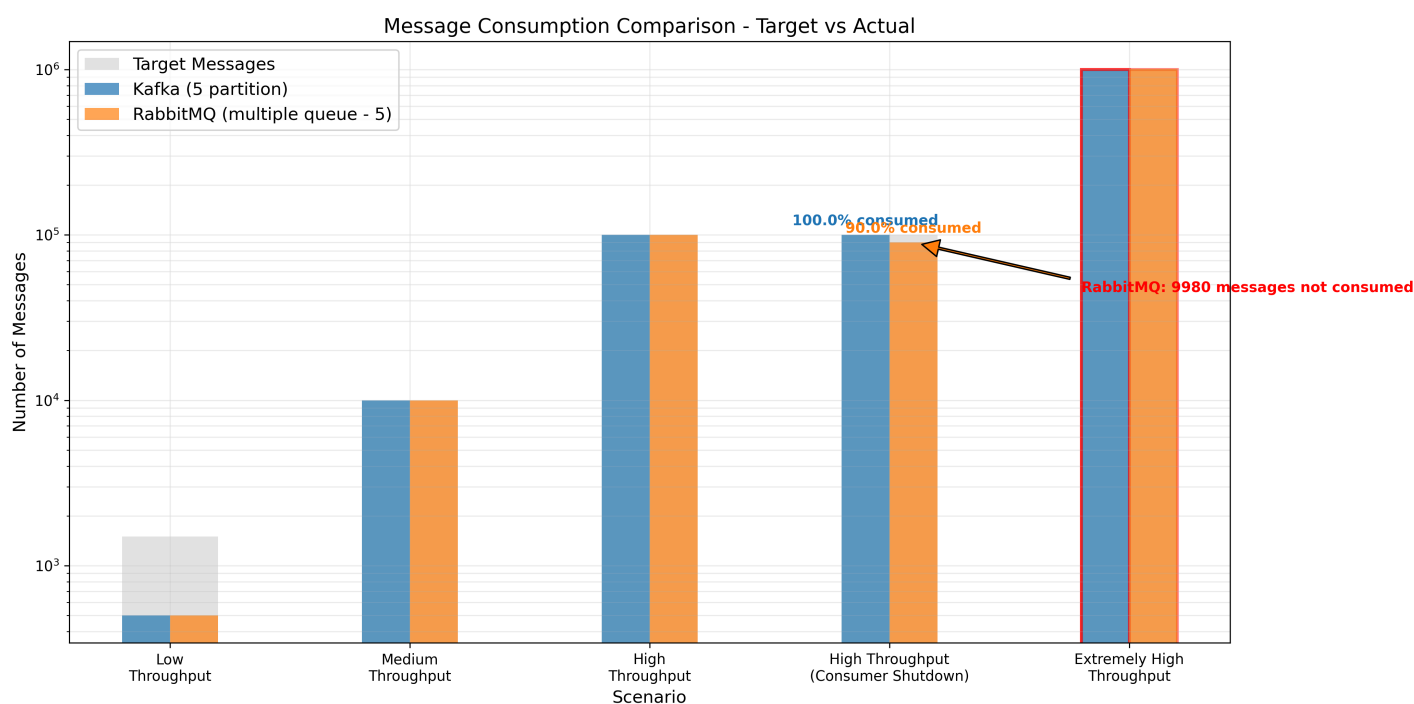


Figure 2: Comparison of target and actual message consumption between Kafka and RabbitMQ. Note the significant difference in the Consumer Disconnect scenario.

The chart illustrates:

- Both systems process 100% of messages in normal scenarios
- In the Consumer Disconnect scenario, Kafka processes all messages while RabbitMQ processes approximately 90%
- The message consumption difference highlights architectural trade-offs between the two systems

Key Findings

From the benchmark results, we observe:

- **Kafka** generally provides lower latency per message than RabbitMQ across all scenarios
- **RabbitMQ** shows higher latency but maintains reliable message delivery

- When a consumer disconnects:
 - Kafka sees a modest increase in latency
 - RabbitMQ experiences a larger increase
 - Both systems continue to process messages reliably

Performance Comparison

Scenario	Kafka Avg Latency (ms)	RabbitMQ Avg Latency (ms)	Kafka E2E Latency (s)	RabbitMQ E2E Latency (s)
Low Throughput	6.18	0.26	3.09	0.13
Medium Throughput	0.44	0.47	4.43	4.71
High Throughput	0.20	0.37	20.20	36.76
Consumer Disconnect	0.78	0.40	77.60	35.93
Extreme Throughput	0.17	0.36	167.54	363.23

Understanding the Consumer Disconnect Scenario

The Consumer Disconnect scenario reveals interesting differences between Kafka and RabbitMQ:

1. Message Processing Count:

- Kafka processed all 100,000 messages despite a consumer disconnection
- RabbitMQ processed 90,020 messages (about 10% message loss)

2. Latency Patterns:

- **Kafka:** Shows significantly higher latency (0.78ms per message, compared to 0.20ms in High Throughput)
- **RabbitMQ:** Maintains similar latency (0.40ms per message, close to its normal 0.37ms)

3. End-to-End Latency:

- **Kafka:** Exhibits much higher E2E latency (77.60s vs. 20.20s in High Throughput)
- **RabbitMQ:** Shows slightly lower E2E latency (35.93s vs. 36.76s)

This unexpected pattern can be explained by key architectural differences:

- **Kafka's Partition Rebalancing:** When a consumer disconnects, Kafka initiates partition rebalancing, which introduces a temporary processing delay but ensures that all messages are eventually processed. This explains the higher latency but complete message processing.
- **RabbitMQ's Queue Isolation:** RabbitMQ's multiple queue approach means that when a consumer disconnects, only the messages assigned to that specific queue remain unprocessed. The remaining consumers continue processing their assigned queues efficiently, resulting in lower overall latency but incomplete message processing.
- **Trade-off Implications:** This demonstrates an important trade-off between the two systems - Kafka prioritizes complete message processing at the cost of temporary performance degradation, while RabbitMQ maintains performance stability at the risk of message loss when consumers disconnect.

Analysis and Conclusions

1. Throughput Scaling:

- Kafka maintains consistently low latency even as throughput increases
- RabbitMQ's latency increases more dramatically under higher loads

2. Failure Resilience:

- Both systems handle consumer failures well
- Kafka recovers slightly faster due to its partition-based architecture
- RabbitMQ has higher latency during recovery but ensures reliable message delivery

3. Use Case Recommendations:

- For high-throughput, low-latency applications: **Kafka**
- For guaranteed delivery with simpler setup: **RabbitMQ**
- For systems requiring consumer failover: Both work well, with Kafka having a slight edge

This benchmark demonstrates that modern message queue systems can handle significant throughput while maintaining reasonable latency, making them suitable for a wide range of distributed application scenarios.